

(No)SQL

Databases

Agenda

- Who am I
- Classification of Databases
- Types of NoSQL
- Under the hood (Distributed systems)
- Benchmarking Tools
- Personal Journey
- Discussions\Q&A

Who am I

- Suyambu Ganesh Thanga Nadar
 - Call sign : **Ganesh**
- Married, two children, living in Amstelveen
- Senior Java Dev/Tech Lead/Solution Architect (Cloud\μServices\IoT)
- Joined Quby in 2016\Merged with Eneco in 2021
- Tech enthusiast\Innovative\fascinated by BigData and IoT
- Prof. career spanning 15+ years



Classification

- SQL - Relational Databases
- NoSQL - Not Only SQL
- NewSQL - NoSQL + SQL

SQL

- Relational Databases
- Based on Relational Algebra not because of Foreign key
- Predefined Schema in Tabular
- ACID Properties
 - Atomicity (Commits entire operation or rollbacks to prev. state)
 - Consistency (Maintains the data integrity, same view for all clients)
 - Isolation (R/W operations are performed in isolation)
 - Durability (Successful commits should be permanent)
- Vertical scalability
- SQL as the query language
- OLTP (Transaction Processing - Tables)
- OLAP (Business Intelligence - Cubes)
- Synchronous Inserts and Updates
- Storage + Query + Index engines

ORACLE®
DATABASE



NoSQL

- Non – relational or Distributed databases
- Dynamic schema
- Document\Key Value\Graph\Wide Column store
- Unstructured query language
- Based on CAP theorem
 - Consistency (Same view for all clients)
 - Availability (Highly available even half of machine is down)
 - Partition tolerant (The system works even when there is partial network failure)
- Not suited for huge Transactional workloads
- Horizontally scalable (115k nodes and over 10PB of data)
- Web 2.0 (user generated content and web)
- Asynchronous inserts and updates




Cloud
Bigtable



amazon
DynamoDB

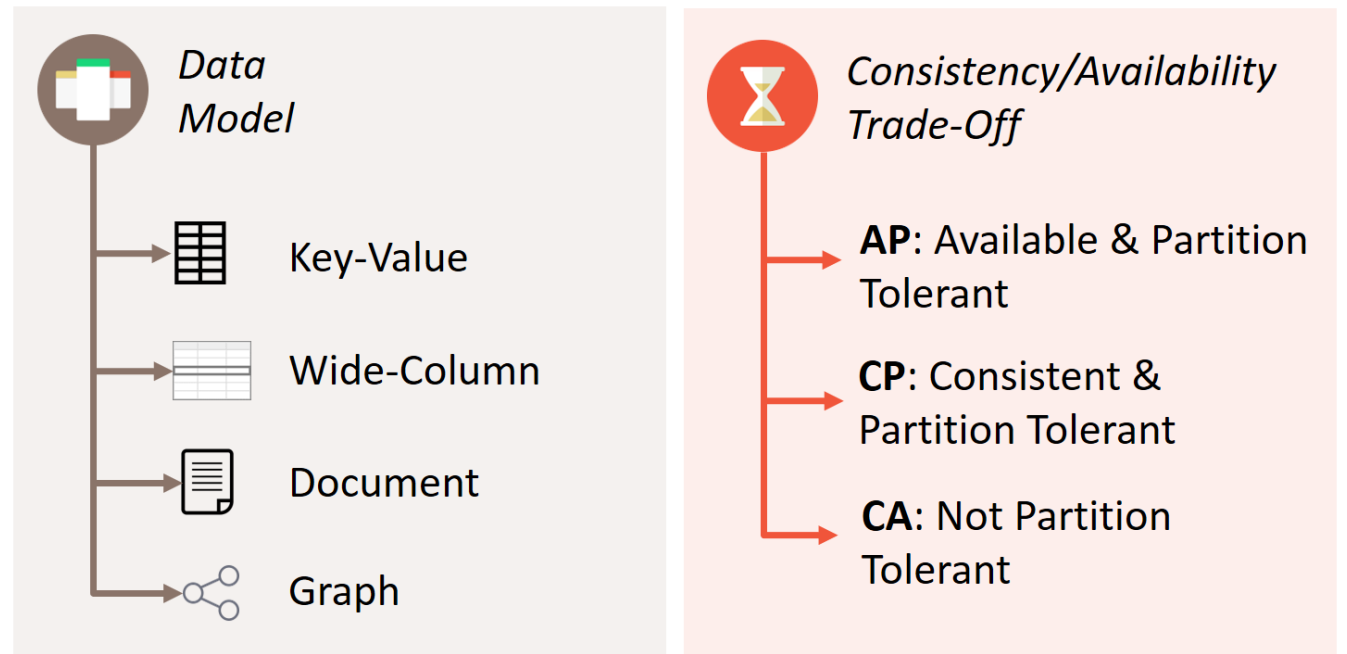
NewSQL

- New Generation of SQL databases
- ACID properties + Distributed Goodness
- Compromise on latency
- SQL and Table structure
- Horizontally scalable
- Globally distributed




Type of NoSQL

- Two approaches
 1. DataModel
 2. CAP theorem (Professor Eric A. Brewer)



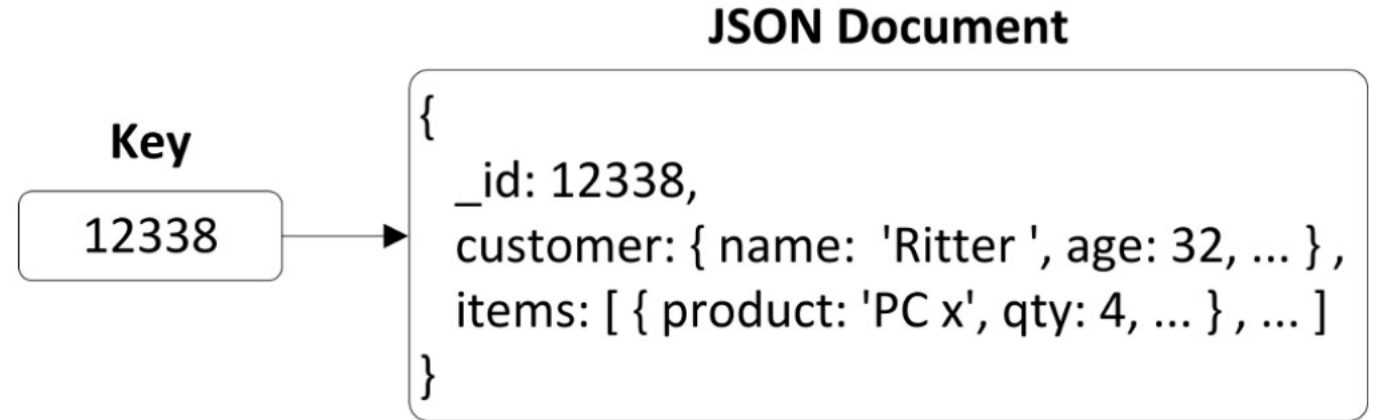
Key value Store

- KV pairs
- Unique Keys
- Schemaless
- Fast read and writes
- Inmemory storage
- Caching
- Counters
- FDB (billions of databases)



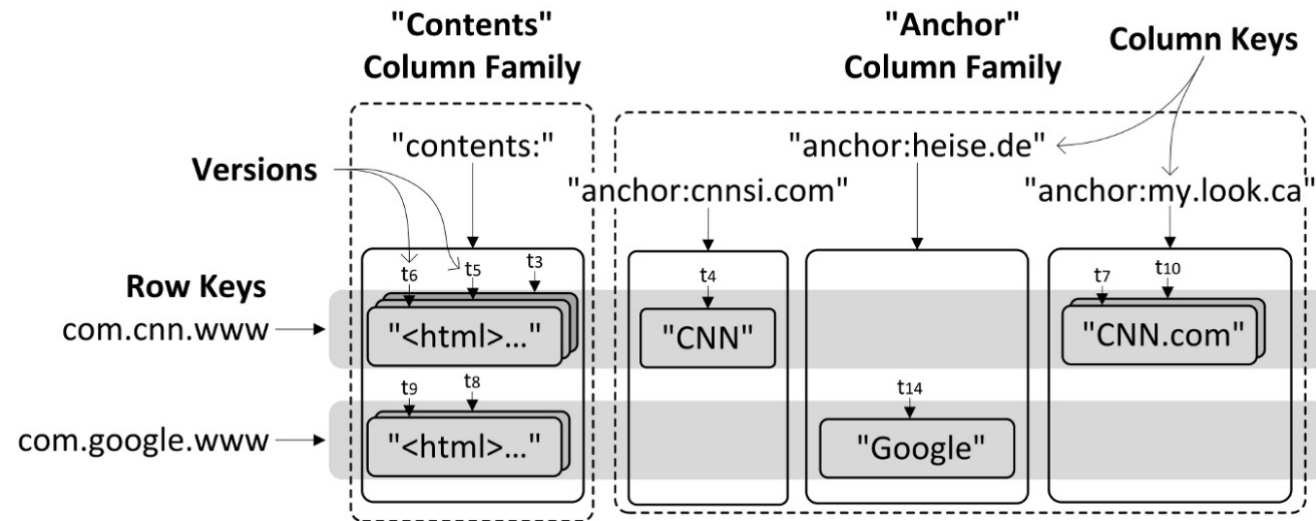
Document Store

- Key Value
- Semi-structured JSON
- Full document or part of Do
- Aggregation
- Full text search
- Analytics
- Product catalog
- Product search




Wide column store

- Similar to Map structure
- Row keys and Column keys
- Column families
- Billions of columns and rows
- Petabyte scale
- 100 millions of write/sec
- Google Bigtable pioneered
- User generated content
- Events\Time series data



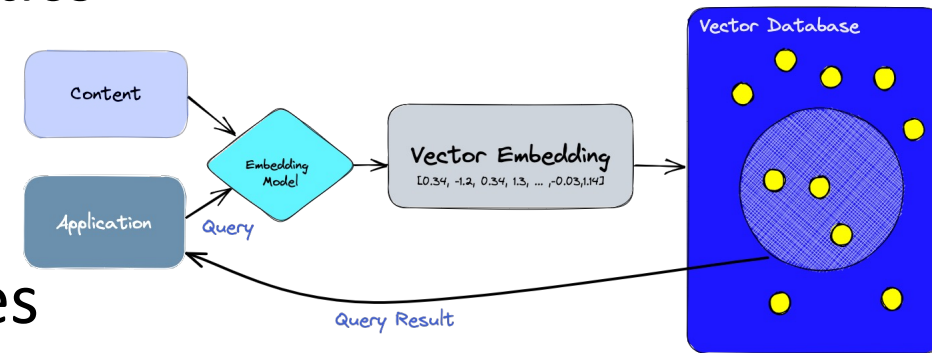
Graph store

- Graph theory
- Nodes, edges and properties
- Efficient join
- Fast traversal
- Millions of connections/second
- Recommendation Engine
- Master data management
- Fraud detection
- Identity and Access management



Vector Databases (HOT Topic right now)

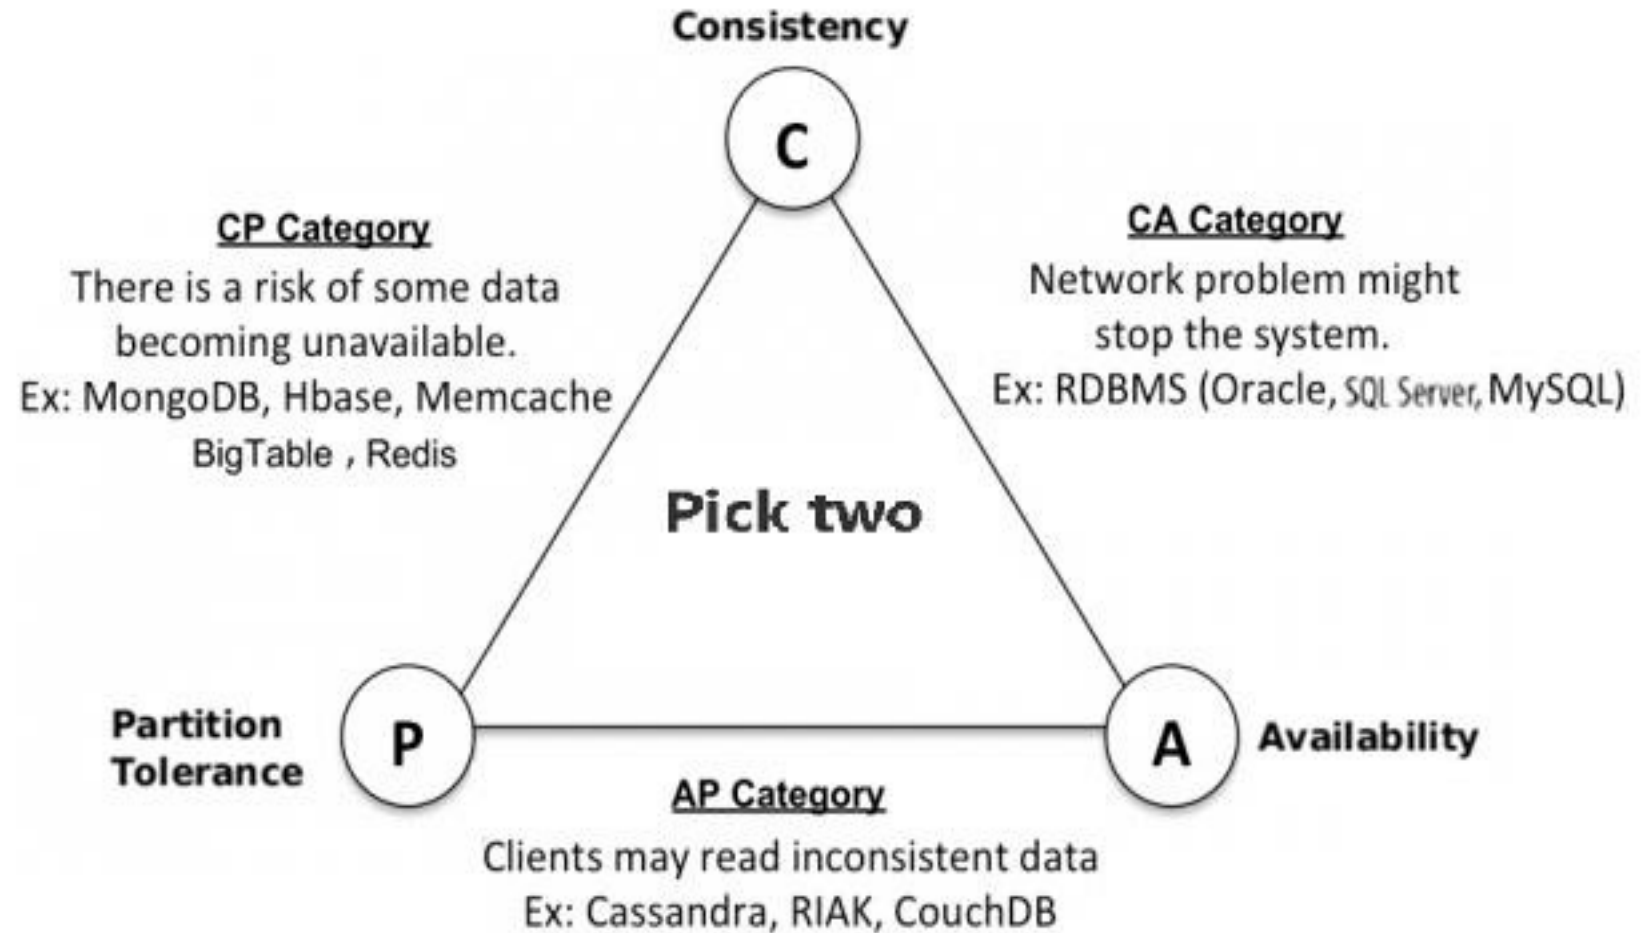
- Stores vector embeddings
- Embeddings are encoded representation of text, images, audio
- Stored as a array of numbers with the semantics
- Vector db supports similarity search
- Used in AI applications like ChatGPT
- Highly scalable, low latency, Real time updates



CAP Theorem

- Defines upper bound guarantees on Distributed systems
 - **Consistency (C):**
 - Reads and writes are always executed atomically and are strictly consistent.
 - All clients have the same view on the data at all times.
 - **Availability (A):**
 - Every non-failing node in the system can always accept read and write requests by clients and will eventually return with a success response.
 - **Partition-tolerance (P):**
 - The system upholds the previously displayed consistency guarantees and availability in the presence of message loss between the nodes or partial system failure.

CAP Theorem



Comparision

	Funct. Req.								Non-Funct. Req.										Techniques																					
	Scan Queries	ACID Transactions	Conditional Writes	Joins	Sorting	Filter Queries	Full-Text Search	Analytics	Data Scalability	Write Scalability	Read Scalability	Elasticity	Consistency	Write Latency	Read Latency	Write Throughput	Read Availability	Write Availability	Durability	Range-Sharding	Hash-Sharding	Entity-Group Sharding	Consistent Hashing	Shared-Disk	Transaction Protocol	Sync. Replication	Async. Replication	Primary Copy	Update Anywhere	Logging	Update-in-Place	Caching	In-Memory Storage	Append-Only Storage	Global Indexing	Local Indexing	Query Planning	Analytics Framework	Materialized Views	
MongoDB	x		x		x	x	x	x	x	x	x		x	x	x	x	x		x	x	x					x	x		x			x	x			x	x	x		
Redis	x	x	x								x		x	x	x	x	x		x								x	x		x			x							
HBase	x		x		x			x	x	x	x	x	x	x		x			x	x					x			x		x			x							
Riak							x	x	x	x	x	x		x	x	x	x	x	x		x		x				x			x	x	x			x	x		x		
Cassandra	x		x		x		x	x	x	x	x	x		x		x	x	x	x		x		x				x		x	x		x		x	x	x				x
MySQL	x	x	x	x	x	x	x	x			x		x						x					x	x		x	x		x	x	x				x	x			

Under the HOOD of NoSQL

- Distributed computing
 - Responsibilities shared across many computers in a network
- Shared nothing architecture
 - Each server instance manages its own data
- Consensus Algorithm
 - leader election
 - data sharding\ data partitioning
 - Metadata sharing
 - example algorithms
 - Paxos (cassandra)
 - Raft (Neo4j, YugabyteDB)
 - ZAB (Zookeeper)

Benchmarking Tools

- YCSB : Yahoo cloud service benchmark
 - <https://github.com/brianfrankcooper/YCSB>
- TLA+
 - Is a spec to design systems and algorithms and verifying
 - <https://learntla.com/introduction/>
- TPC
 - Transaction Processing Performance Council
 - <http://www.tpc.org/>
- Jepsen
 - To improve the safety of the data.
 - Verify databases on ACID or CAP compliance
 - <https://jepsen.io/>

Personal Journey with NoSQL databases

- Started in 2012
 - Designed logging and tracing system for KLM
 - Storing and Indexing 70 GB of XML data in 4 minutes
 - Used Cassandra 0.8 in 3 un-used desktop machines
 - 1 month data retention of 6 TB of compressed XML
 - Searching trace by PNR, TicketNumber in < 30ms
- In 2013
 - Evaluated MongoDB, Cassandra and CouchBase
 - Recommended Couchbase to Hotel Chain Marriott
 - Still couchbase is in use at Marriott



[← Back to customer page](#)

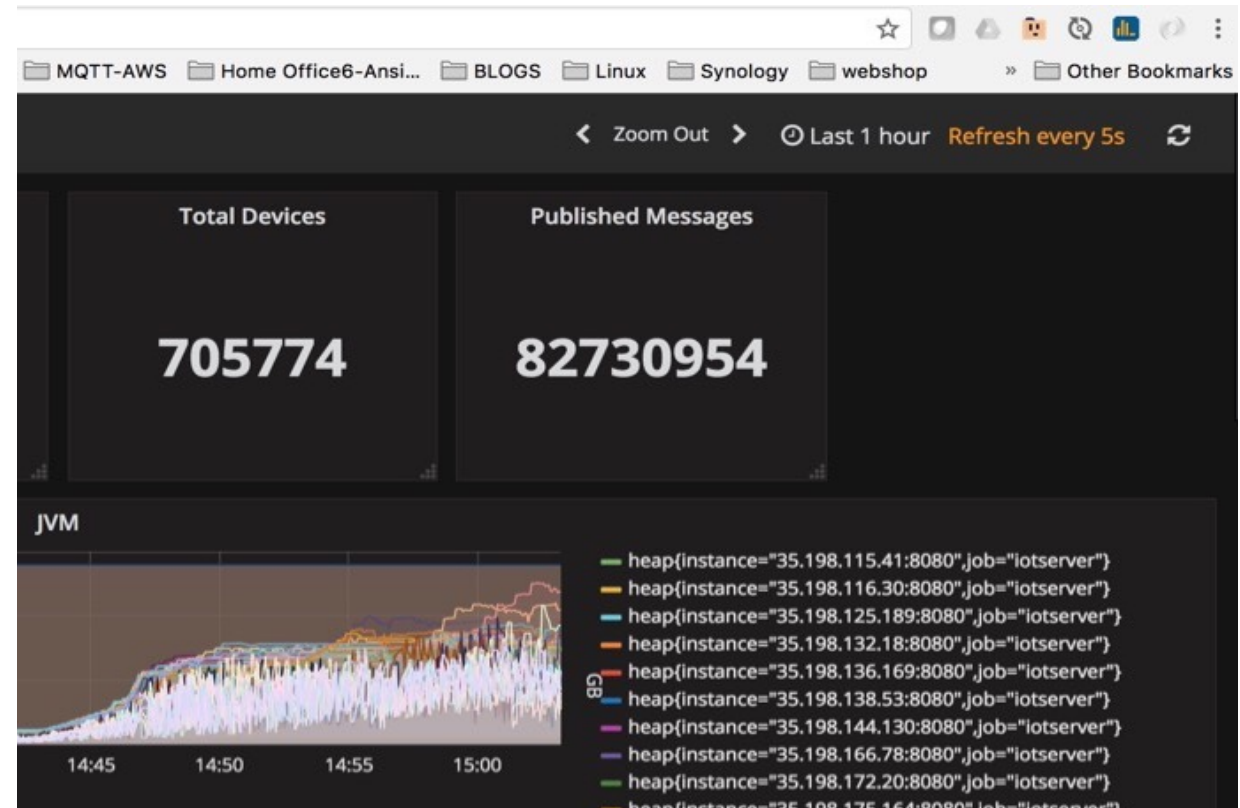


Marriott - Creating a modern customer experience for the digital age

To maintain its competitive edge in the digital economy, Marriott wanted to create personalized customer experiences, improve online reliability, and release new apps faster. After an evaluation against MongoDB™ and Cassandra, Marriott chose Couchbase to replace its legacy infrastructure. Couchbase had already proven itself in the industry, and Marriott's solutions architects were impressed by Couchbase's built-in cache, ease and flexibility for moving and adding cluster nodes, and easy disaster recovery.

Personal Journey with NoSQL databases

- Created MQTT Broker
- Cassandra as backend
- Simulated 700K MQTT clients
- 82 Million messages in 15 minutes
- 90K msgs/second



Q/A